

Operating Systems

Isfahan University of Technology
Electrical and Computer Engineering Department

Zeinab Zali

Synchronization-Condition Variables-monitors

Condition Variable

- there are many cases where a thread wishes to check whether a condition is true before continuing its execution
- To wait for a condition to become true, a thread can make use of what is known as a condition variable
- A condition variable is an **explicit queue** that threads can put themselves on when some state of execution (i.e., some condition) is not as desired (by waiting on the condition)
- some other thread, when it changes said state, can then **wake one (or more)** of those waiting threads and thus allow them to continue (by signaling on the condition)


CV definitions and routines

- `pthread_cond_t c`
- `pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m):`
The `wait()` call is executed when a thread wishes to put itself to sleep
 - The responsibility of `wait()` is to release the lock and put the calling thread to sleep (atomically); when the thread wakes up (after some other thread has signaled it), it must re-acquire the lock before returning to the caller.
- `pthread_cond_signal(pthread_cond_t *c):` the `signal()` call is executed when a thread has changed something in the program and thus wants to wake a sleeping thread waiting on this condition

CV example 1

```
1 void *child(void *arg) {
2     printf("child\n");
3     // XXX how to indicate we are done?
4     return NULL;
5 }
6
7 int main(int argc, char *argv[]) {
8     printf("parent: begin\n");
9     pthread_t c;
10    Pthread_create(&c, NULL, child, NULL); // create child
11    // XXX how to wait for child?
12    printf("parent: end\n");
13    return 0;
14 }
```

Desired	parent: begin
Output:	child
	parent: end

```
1  int done  = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c  = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6
7
8
9
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14      thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19
20
21
22
23 }
24
25 int main(int argc, char *argv[]) {
26     printf("parent: begin\n");
27     pthread_t p;
28     Pthread_create(&p, NULL, child, NULL);
29     thr_join();
30     printf("parent: end\n");
31     return 0;
32 }
```

What is the problem with this solution?

```
1  void thr_exit() {
2      Pthread_mutex_lock(&m);
3      done = 1;
4      Pthread_cond_signal(&c);
5      Pthread_mutex_unlock(&m);
6  }
7  void thr_join() {
8      Pthread_mutex_lock(&m);
9      while (done == 0)
10         Pthread_cond_wait(&c, &m);
11         Pthread_mutex_unlock(&m);
12     }
```

What is the problem with this solution?

```
1  void thr_exit() {
2      done = 1;
3      Pthread_cond_signal(&c);
4  }
5
6  void thr_join() {
7      if (done == 0)
8          Pthread_cond_wait(&c);
9  }
```

Producer/Consumer Problem

```
1  int loops; // must initialize somewhere...
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8
9
10
11
12
13
14      }
15  }
16
17  void *consumer(void *arg) {
18      int i;
19      for (i = 0; i < loops; i++) {
20
21
22
23
24
25
26
27      }
28  }
```

// p1
// p2
// p3
// p4
// p5
// p6

// c1
// c2
// c3
// c4
// c5
// c6

Producer/Consumer Problem

```
1  int buffer;  
2  int count = 0; // initially, empty  
3  
4  void put(int value) {  
5      assert(count == 0);  
6      count = 1;  
7      buffer = value;  
8  }  
9  
10 int get() {  
11     assert(count == 1);  
12     count = 0;  
13     return buffer;  
14 }
```

Producer/Consumer Problem

```
1  int loops; // must initialize somewhere...
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);           // p1
9          if (count == 1)                       // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                             // p4
12             Pthread_cond_signal(&cond);         // p5
13             Pthread_mutex_unlock(&mutex);       // p6
14         }
15     }
16
17     void *consumer(void *arg) {
18         int i;
19         for (i = 0; i < loops; i++) {
20             Pthread_mutex_lock(&mutex);         // c1
21             if (count == 0)                     // c2
22                 Pthread_cond_wait(&cond, &mutex); // c3
23             int tmp = get();                     // c4
24             Pthread_cond_signal(&cond);         // c5
25             Pthread_mutex_unlock(&mutex);       // c6
26             printf("%d\n", tmp);
27         }
28     }
```

Producer/Consumer broken solution with if

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	T_{c1} awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	T_{c2} sneaks in ...
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	T_p awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

```

1  int loops;
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);          // p1
9          while (count == 1)                  // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                          // p4
12             Pthread_cond_signal(&cond);      // p5
13             Pthread_mutex_unlock(&mutex);    // p6
14         }
15     }
16
17     void *consumer(void *arg) {
18         int i;
19         for (i = 0; i < loops; i++) {
20             Pthread_mutex_lock(&mutex);      // c1
21             while (count == 0)                // c2
22                 Pthread_cond_wait(&cond, &mutex); // c3
23             int tmp = get();                  // c4
24             Pthread_cond_signal(&cond);      // c5
25             Pthread_mutex_unlock(&mutex);    // c6
26             printf("%d\n", tmp);
27         }
28     }

```

```

1  int loops;
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);          // p1
9          while (count == 1)                    // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11         ...

```

TIP: USE WHILE (NOT IF) FOR CONDITIONS

When checking for a condition in a multi-threaded program, using a `while` loop is always correct; using an `if` statement only might be, depending on the semantics of signaling. Thus, always use `while` and your code will behave as expected.

```

18      int i;
19      for (i = 0; i < loops; i++) {
20          Pthread_mutex_lock(&mutex);          // c1
21          while (count == 0)                    // c2
22             Pthread_cond_wait(&cond, &mutex); // c3
23          int tmp = get();                      // c4
24          Pthread_cond_signal(&cond);          // c5
25          Pthread_mutex_unlock(&mutex);        // c6
26          printf("%d\n", tmp);
27      }
28  }

```

Producer/Consumer still a broken solution

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	T_{c1} awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	T_{c1} grabs data
c5	Running		Ready		Sleep	0	Oops! Woke T_{c2}
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	Everyone asleep...

Producer/Consumer with buffer size MAX

```
1  cond_t empty, fill;
2  mutex_t mutex;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          Pthread_mutex_lock(&mutex);          // p1
8          while (count == MAX)                 // p2
9              Pthread_cond_wait(&empty, &mutex); // p3
10         put(i);                               // p4
11         Pthread_cond_signal(&fill);           // p5
12         Pthread_mutex_unlock(&mutex);         // p6
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         Pthread_mutex_lock(&mutex);          // c1
20         while (count == 0)                   // c2
21             Pthread_cond_wait(&fill, &mutex); // c3
22         int tmp = get();                      // c4
23         Pthread_cond_signal(&empty);         // c5
24         Pthread_mutex_unlock(&mutex);        // c6
25         printf("%d\n", tmp);
26     }
27 }
```

Producer/Consumer with buffer size MAX

```
1  int buffer[MAX];
2  int fill_ptr = 0;
3  int use_ptr  = 0;
4  int count    = 0;
5
6  void put(int value) {
7      buffer[fill_ptr] = value;
8      fill_ptr = (fill_ptr + 1) % MAX;
9      count++;
10 }
11
12 int get() {
13     int tmp = buffer[use_ptr];
14     use_ptr = (use_ptr + 1) % MAX;
15     count--;
16     return tmp;
17 }
```




Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- Pseudocode syntax of a monitor:

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure P2 (...) { ... }


    procedure Pn (...) {.....}

    initialization code (...) { ... }
}
```





Schematic view of a Monitor





Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex  
mutex = 1
```

- Each procedure **P** is replaced by

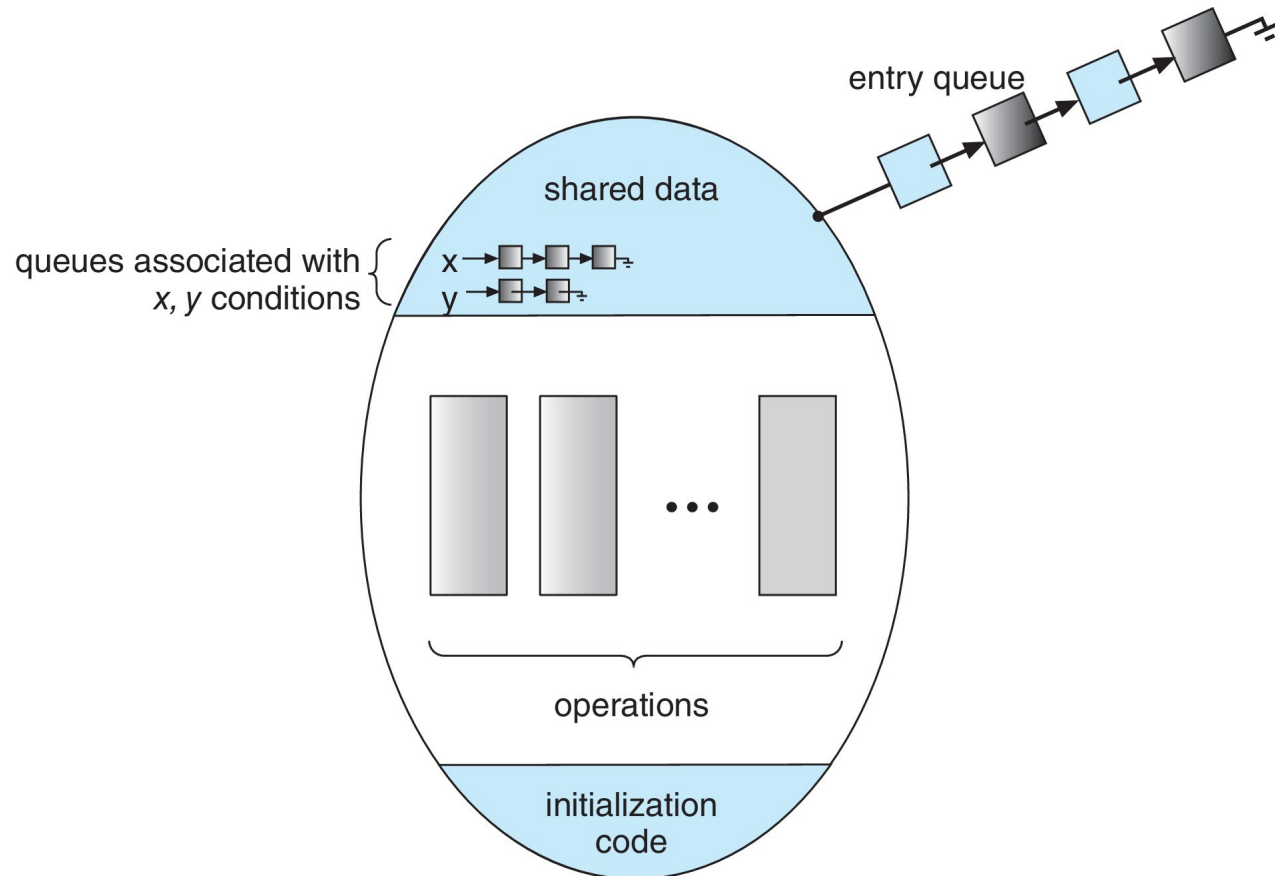
```
wait(mutex) ;  
    ...  
    body of P ;  
    ...  
signal(mutex) ;
```

- Mutual exclusion within a monitor is ensured





Monitor with Condition Variables





Usage of Condition Variable Example

- Consider P_1 and P_2 that need to execute two statements S_1 and S_2 and the requirement that S_1 to happen before S_2

- Create a monitor with two procedures F_1 and F_2 that are invoked by P_1 and P_2 respectively
- One condition variable “x” initialized to 0
- One Boolean variable “done”

- **F1:**

```
 $S_1;$   
  
done = true;  
  
x.signal();
```

- **F2:**

```
if done = false  
    x.wait()  
  
 $S_2;$ 
```





Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
```

```
{
```

```
    enum { THINKING, HUNGRY, EATING} state [5] ;
```

```
    condition self [5];
```

```
    void pickup (int i) {  
        }
```

```
    void putdown (int i) {  
    }
```

- Each philosopher “i” invokes the operations **pickup()** and **putdown()** in the following sequence:

```
DiningPhilosophers.pickup(i) ;
```

```
    /** EAT **/
```

```
DiningPhilosophers.putdown(i) ;
```





Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING, HUNGRY, EATING} state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```





Solution to Dining Philosophers (Cont.)

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```





Solution to Dining Philosophers (Cont.)

- Each philosopher “i” invokes the operations `pickup()` and `putdown()` in the following sequence:

```
DiningPhilosophers.pickup(i);
```

```
/** EAT **/
```

```
DiningPhilosophers.putdown(i);
```

- No deadlock, but starvation is possible





End of Chapter 6

